

Deep Learning

Lab 5: Regularization

Chia-Hung Yuan & DataLab

Department of Computer Science,
National Tsing Hua University, Taiwan

Regularization

- Regularization refers to techniques that improve the generalizability of a trained model

Outline

- Scikit-learn
- Learning Theory
 - Error Curves and Model Complexity
 - Learning Curves and Sample Complexity
- Weight Decay
 - Ridge Regression
 - LASSO
- Validation
- Assignment

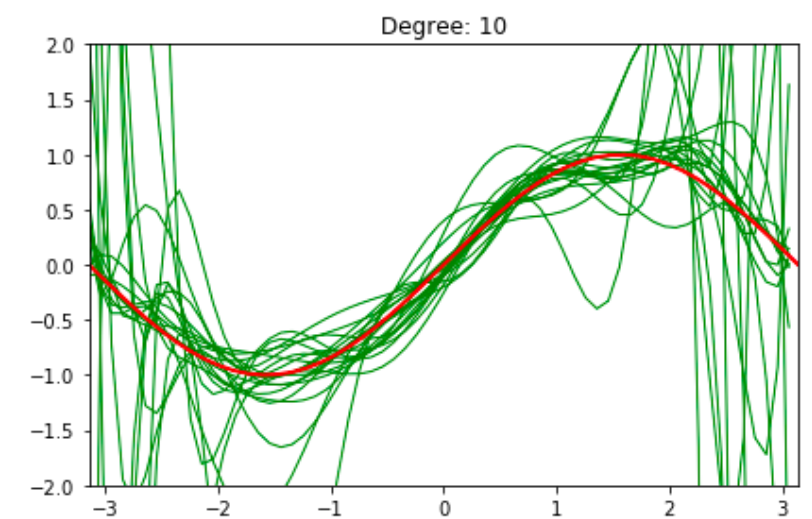
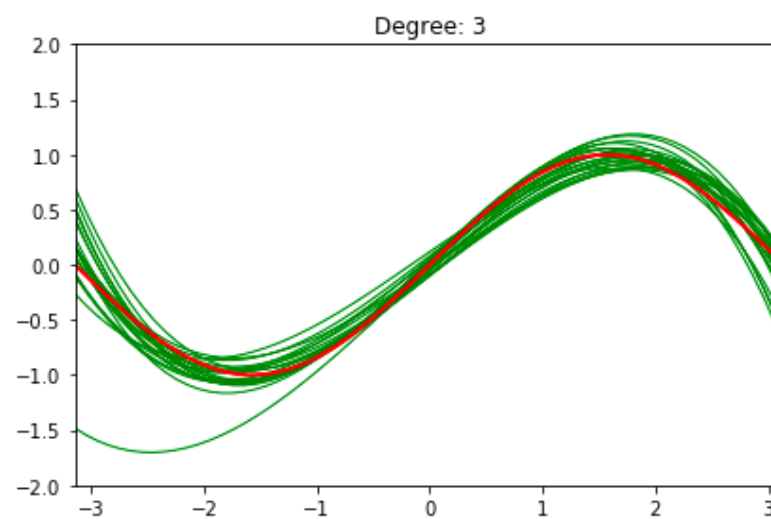
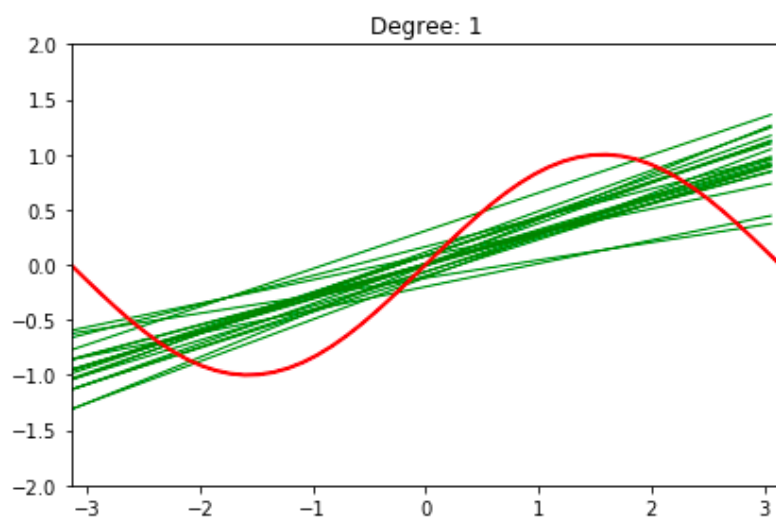
Scikit-learn

- Scikit-learn is a free software machine learning library for the Python programming language
- It features various classification, regression and clustering algorithms
 - including support vector machines, random forests, gradient boosting, k-means and DBSCAN, and is designed to interoperate with the Python numerical and scientific libraries NumPy and SciPy
- `pip install scikit-learn` / `conda install scikit-learn`



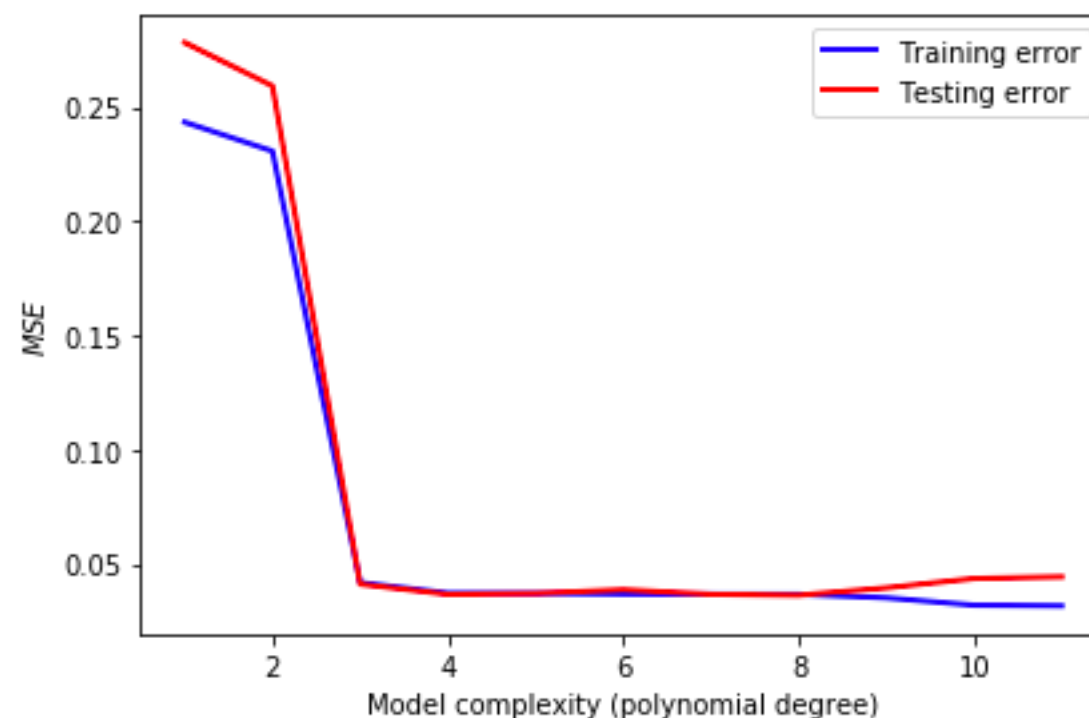
Learning Theory

- Learning theory provides a means to understand the generalizability of the model
- **Model complexity** plays a crucial role
 - Too simple: high bias and underfitting
 - Too complex: high variance and overfitting



Error Curves and Model Complexity

- It is relatively hard to observe the figures showed in the last slide, since normally we will never know the data distribution of ground truth
- Instead, we can get those information by observing the training and testing error

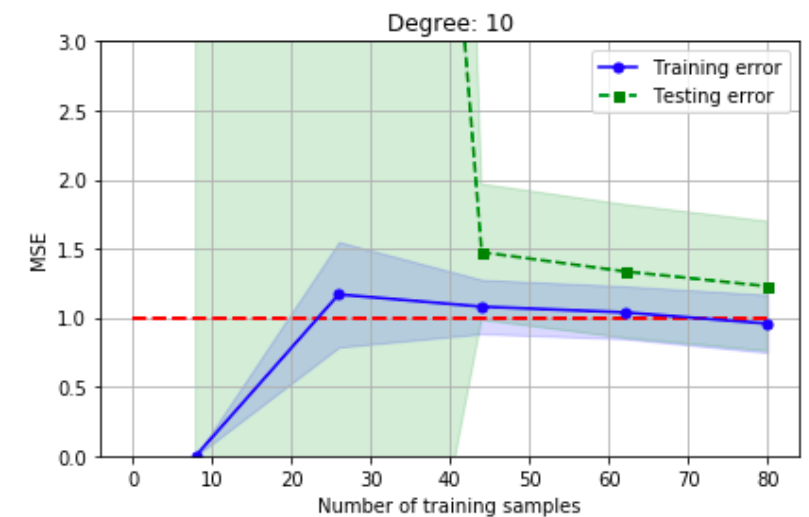
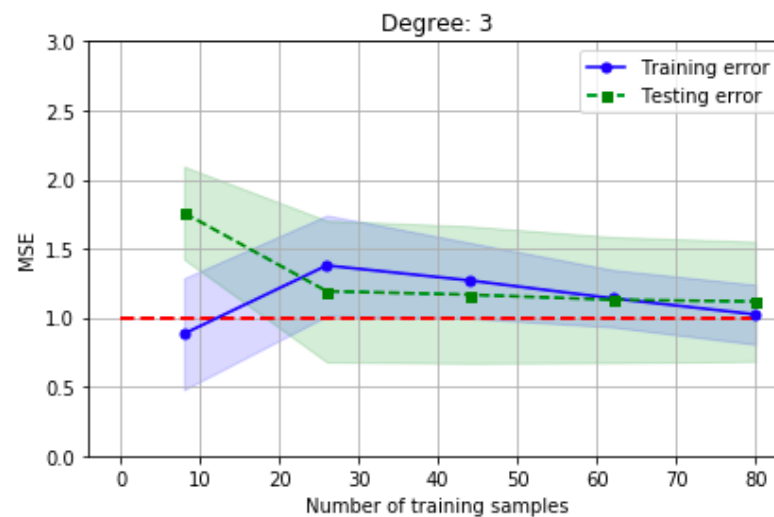
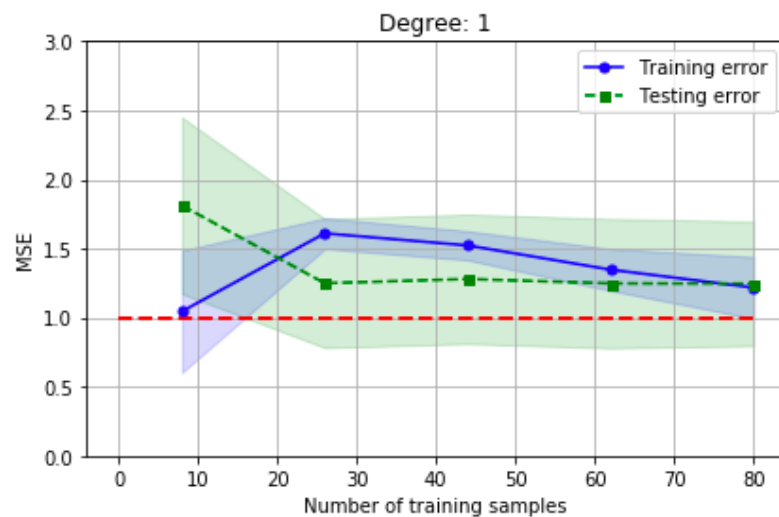


Error Curves and Model Complexity

- Although the error curve visualizes the impact of model complexity, the bias-variance tradeoff holds only when you have sufficient training examples

Learning Curves and Sample Complexity

- The bounding methods of learning theory tell us that a model is likely to overfit regardless of its complexity **when the size of training set is small**. The **learning curves** are a useful tool for understanding how much training examples are sufficient



Weight Decay

- A common regularization approach. The idea is to add a term in the cost function against complexity

- Ridge Regression (L_2)

$$\arg \min_{\mathbf{w}, b} \|\mathbf{y} - (X\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|^2$$

- LASSO (L_1)

$$\arg \min_{\mathbf{w}, b} \|\mathbf{y} - (X\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|_1$$

Ridge Regression

- A small value α drastically reduces the testing error. Nevertheless, it's not a good idea to increase α forever, since it will over-shrink the coefficients of w and result in underfitting

$$\arg \min_{w,b} \|y - (Xw - b\mathbf{1})\|^2 + \alpha \|w\|^2$$

```
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

poly = PolynomialFeatures(degree=3)
X_poly = poly.fit_transform(X_std)
X_train, X_test, y_train, y_test = train_test_split(
    X_poly, y, test_size=0.3, random_state=0)

for a in [0, 1, 10, 100, 1000]:
    lr_rg = Ridge(alpha=a)
    lr_rg.fit(X_train, y_train)

    y_train_pred = lr_rg.predict(X_train)
    y_test_pred = lr_rg.predict(X_test)

    print('\n[Alpha = %d]' % a )
    print('MSE train: %.2f, test: %.2f' % (
        mean_squared_error(y_train, y_train_pred),
        mean_squared_error(y_test, y_test_pred)))
```

```
[Alpha = 0]
MSE train: 0.00, test: 19958.68

[Alpha = 1]
MSE train: 0.73, test: 23.05

[Alpha = 10]
MSE train: 1.66, test: 16.83

[Alpha = 100]
MSE train: 3.60, test: 15.16

[Alpha = 1000]
MSE train: 8.81, test: 19.22
```

LASSO

- An alternative weight decay approach that can lead to sparse w is the LASSO. Depending on the value of α , certain weights can become zero much faster than others

$$\arg \min_{w,b} \|y - (Xw - b\mathbf{1})\|^2 + \alpha \|w\|_1$$

```
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error

for a in [0.001, 0.01, 0.1, 1, 10]:
    lr_rg = Lasso(alpha=a)
    lr_rg.fit(X_train, y_train)

    y_train_pred = lr_rg.predict(X_train)
    y_test_pred = lr_rg.predict(X_test)

    print('\n[Alpha = %.2f]' % a )
    print('MSE train: %.2f, test: %.2f' % (
        mean_squared_error(y_train, y_train_pred),
        mean_squared_error(y_test, y_test_pred)))
```

```
[Alpha = 0.0000]
MSE train: 0.55, test: 61.02
```

```
[Alpha = 0.0010]
MSE train: 0.64, test: 29.11
```

```
[Alpha = 0.0100]
MSE train: 1.52, test: 19.51
```

```
[Alpha = 0.1000]
MSE train: 4.34, test: 15.52
```

```
[Alpha = 1.0000]
MSE train: 14.33, test: 22.42
```

```
[Alpha = 10.0000]
MSE train: 55.79, test: 53.42
```

Ridge vs LASSO

- Why is LASSO sparse?

- Ridge Regression (L_2)

$$\arg \min_{\mathbf{w}, b} \|\mathbf{y} - (X\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|^2$$

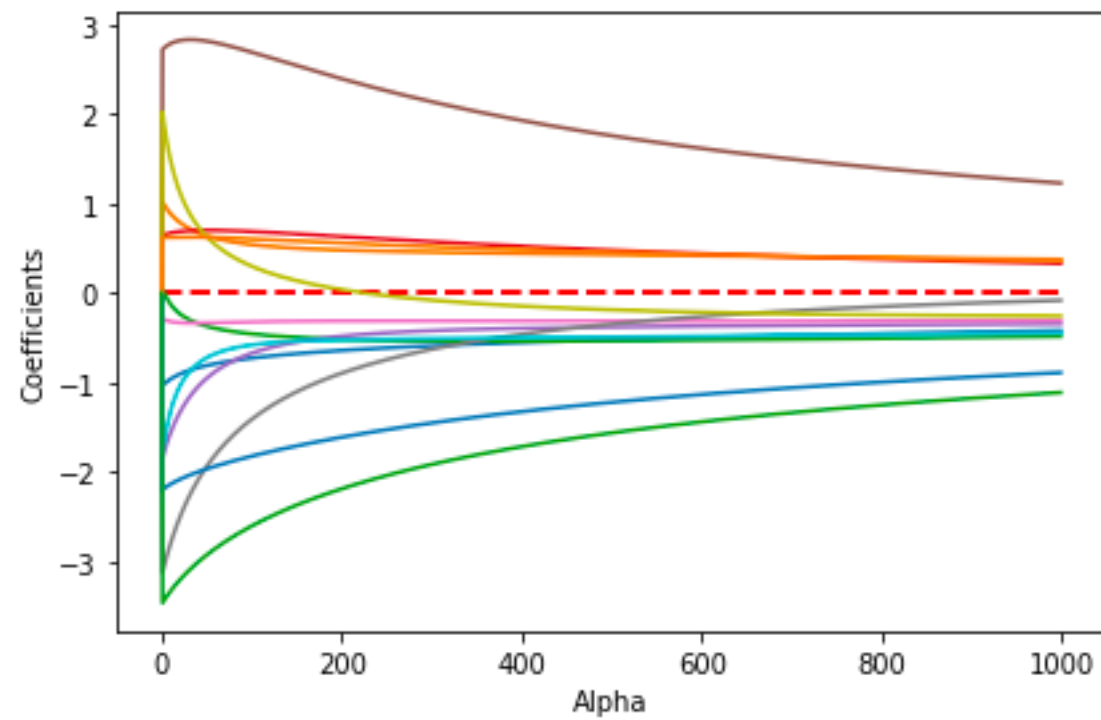
- LASSO (L_1)

$$\arg \min_{\mathbf{w}, b} \|\mathbf{y} - (X\mathbf{w} - b\mathbf{1})\|^2 + \alpha \|\mathbf{w}\|_1$$

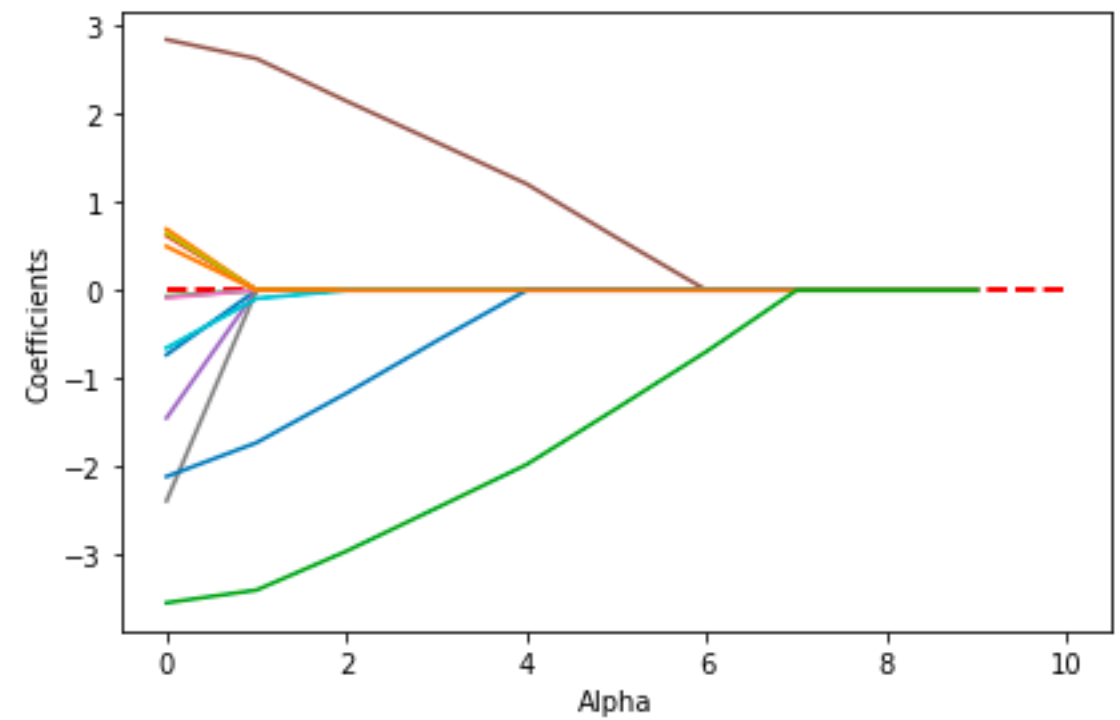


Ridge vs LASSO

Ridge

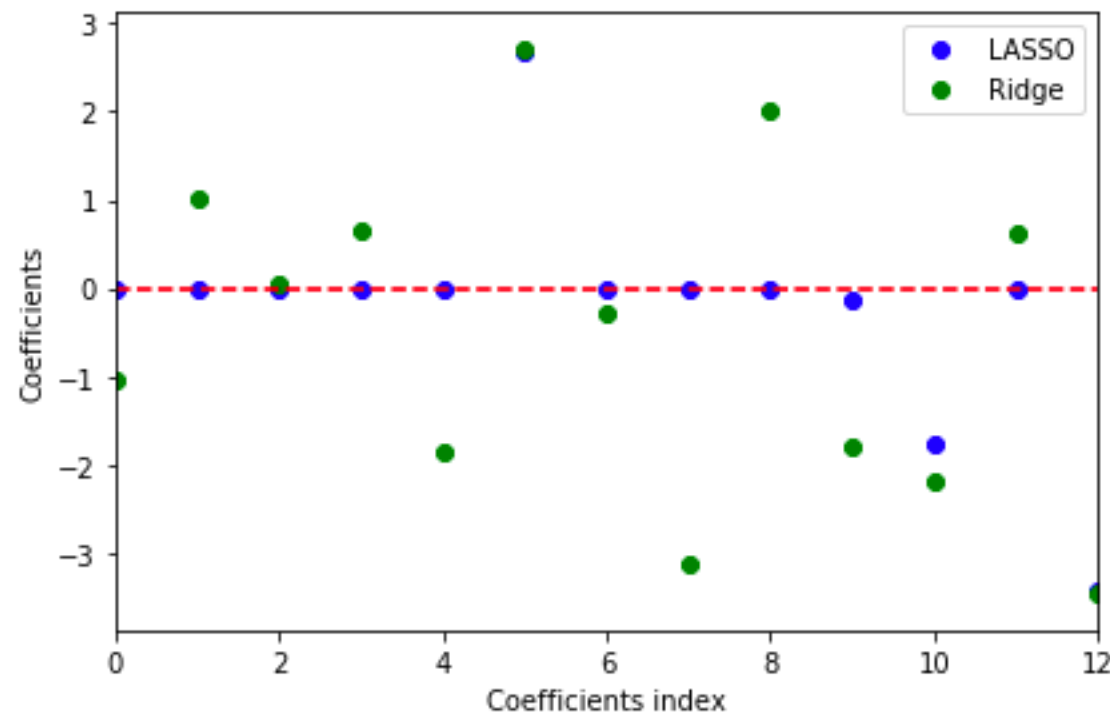


LASSO



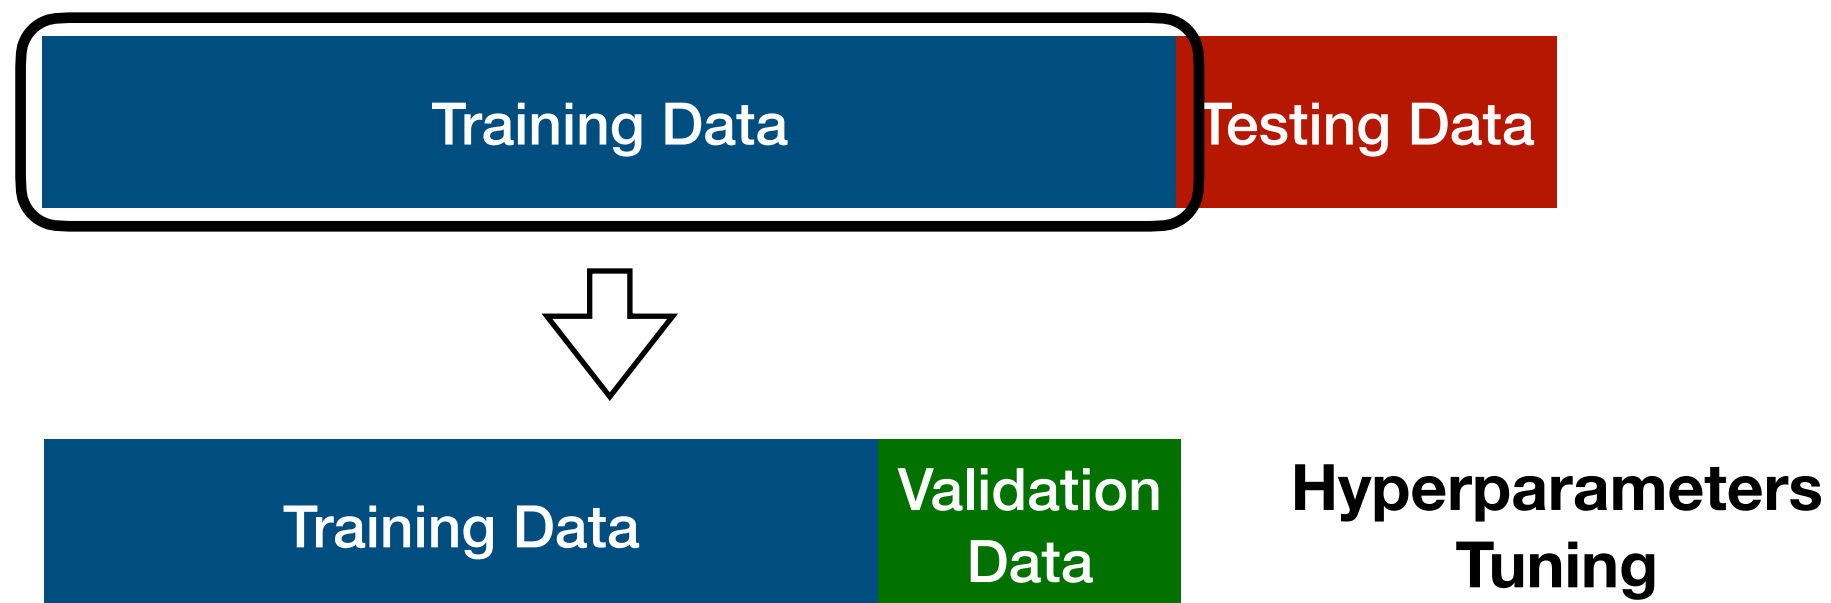
Ridge vs LASSO

- LASSO can also be treated as a supervised feature selection technique when choosing a suitable regularization strength α to make only part of coefficients become exactly zeros



Validation

- Another useful regularization technique that helps us decide the proper value of hyperparameters
- The idea is to split your data into the training, validation, and testing sets and then select the best value based on validation performance
- NOTE: It is important that we should never peek testing data during training



Validation

[Degree = 1]
MSE train: 25.00, valid: 21.43, test: 32.09

[Degree = 2]
MSE train: 9.68, valid: 14.24, test: 20.24

[Degree = 3]
MSE train: 3.38, valid: 17.74, test: 18.63



[Degree = 4]
MSE train: 1.72, valid: 16.67, test: 30.98

[Degree = 5]
MSE train: 0.97, valid: 59.73, test: 57.02

[Degree = 6]
MSE train: 0.60, valid: 1444.08, test: 33189.41

Assignment

- In this assignment, we would like to predict the success of shots made by basketball players in the NBA



58.5	43.2	39.0	36.3	54.7	55.3	41.2	40.5	43.0	54.5
55.3	40.7	39.1	40.8	60.7	62.3	41.3	38.1	39.5	57.2
50.9	39.9	40.3	36.9	47.9	48.7	36.7	39.3	40.8	53.8
55.9	38.8	40.6	38.5	41.6	41.2	38.4	40.1	40.5	52.4
52.2	53.3	41.4	39.2	46.0	42.3	42.5	38.4	52.9	50.2
38.7	50.4	50.6	43.8	40.6	39.8	42.6	51.6	47.7	38.9
35.7	30.8	47.8	51.1	55.0	51.7	49.2	47.1	25.2	26.9
10.0	32.9	41.0	40.3	41.7	34.0	25.0	30.5	17.9	10.0
0.0	10.0	5.9	18.8	41.2	10.7	30.0	29.4	10.7	9.1
30.0	16.7	35.7	7.7	31.6	17.4	13.3	23.1	13.2	15.0

Assignment

- In this assignment, we would like to predict the success of shots made by basketball players in the NBA
 - **y_test** is hidden this time
 - Allow to use **any model** you have learned before to achieve the best accuracy
 - Select the best **3 features**, and show the accuracy with only those
- Hint
 - **Preprocess the data** to help your training
 - Since you don't have y_test this time, you may need to **split a validation set** for checking your performance
 - It is possible to use regression model as a classifier, for example [RidgeClassifier](#)

Assignment

- Submit to iLMS with your **ipynb** (Lab05_{student_id}.ipynb) and **y_pred.csv**
- The notebook should contain
 - How you **evaluate** your model
 - **All models** you have tried and the result
 - Plot the **error curve** of your best model and tell if it is **over-fit or not**
 - The top-3 features you find and how you find it
 - A **brief report** what you do in this assignment
- Deadline: **2020-10-08(Thur) 23:59**

Reference

- [Stanford CS229 Machine Learning](#)
- [NBA shot logs](#)

About Competition

- Students will group (2~4 people a group). This class requires **each group of students to prepare a GPU card** to perform the necessary computing. You can follow [this link](#) to decide which GPU card to go for. **NO GPU CARD PROVIDED IN THE CLASS.**
- Sign up [here](#) for your group before 10/13(Tue)

About Competition

InClass Prediction Competition

DataLabCup: Object Detection

Competition for CS565600 Deep Learning

69 teams · 10 months ago

Overview

Data

Notebooks

Discussion

Leaderboard

Rules

Public Leaderboard

Private Leaderboard

This leaderboard is calculated with approximately 50% of the test data.











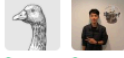




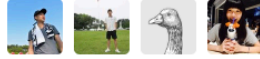
The final results will be based on the other 50%, so the final standings may be different.

📄 Raw Data

🔄 Refresh

#	Team Name	Notebook	Team Members	Score ?	Entries	Last
1	重填一次真的很麻煩		<div><div></div><div></div></div>	0.07996	19	10mo
2	autoencoder		<div><div></div><div></div><div></div></div>	0.23086	28	10mo
3	labXXX		<div><div></div><div></div><div></div><div></div></div>	0.27655	55	10mo
4	acfun02		<div><div></div></div>	0.30113	7	10mo
5	Encoder		<div><div></div><div></div><div></div></div>	0.33275	23	10mo

About Competition

Overview Data Notebooks Discussion <u>Leaderboard</u> Rules							
#	Team Name	Notebook	Team Members	Score ?	Entries	Last	
1	重填一次真的很麻煩			0.07996	19	10mo	
2	autoencoder			0.23086	28	10mo	
3	labXXX			0.27655	55	10mo	
4	acfun02			0.30113	7	10mo	
5	Encoder			0.33275	23	10mo	
6	藝術的狀態			0.33983	22	10mo	
7	Dondon231			0.34905	30	10mo	
	Benchmark-80			0.35104			
8	ChiHang			0.36078	8	10mo	
9	QWQ			0.36667	14	10mo	
10	Benchmark_70			0.41273	15	10mo	
11	Overfitting			0.42421	11	10mo	
12	Human Predict			0.43141	33	10mo	
	Benchmark-60			0.43770			
13	煎魚週			0.45608	7	10mo	
14	华农兄弟			0.50766	34	10mo	